Case Study: Manufacturing IoT Assistant

By Syed Armghan Ahmad – syedarmghanahmad.work@gmail.com

Linkedin - GitHub - Repo for this Project

Overview

As a self-taught developer passionate about IoT and AI, I developed the **Manufacturing IoT Assistant**, a real-time monitoring dashboard and conversational AI tool for factory floor operations. Designed for engineers and plant managers, this application integrates IoT sensor data with AI-driven insights, enabling proactive maintenance and troubleshooting. Built with Python, Streamlit, LlamaIndex, and GroqLLM, the project showcases my ability to combine real-time data processing, interactive visualizations, and natural language processing into a cohesive solution, all learned through self-directed study.

Problem Statement

Modern manufacturing facilities rely on IoT sensors to monitor machine health, but raw sensor data (e.g., vibration, temperature) is often overwhelming and difficult to interpret without specialized tools. Plant managers need real-time insights, predictive maintenance alerts, and actionable guidance to prevent downtime. My goal was to create a tool that:

- Provides a real-time dashboard for monitoring machine status with clear visualizations.
- Offers AI-driven troubleshooting and maintenance recommendations via natural language queries.
- Supports simulation for testing without hardware, making it accessible for development and evaluation.
- Enhances decision-making in high-stakes manufacturing environments.

Approach

I designed a modular, full-stack application that combines a Streamlit-based dashboard with an AI-powered assistant. The system fetches real-time and historical data from ThingSpeak's IoT platform, processes it with LlamaIndex for querying, and visualizes it using Plotly, all while maintaining a user-friendly, dark-mode interface optimized for control rooms.

Key Features

- **Interactive Dashboard**: Displays real-time machine status cards with color-coded alerts, animated gauge charts for vibration and temperature, and historical trend visualizations with configurable time windows.
- **AI Assistant**: Enables natural language queries about machine status, provides maintenance recommendations based on thresholds (e.g., vibration <5 mm/s, temperature <80°C, maintenance every 200 hours), and offers troubleshooting guidance.
- **Simulation Mode**: Generates synthetic sensor data via ThingSpeak's API, allowing testing without physical hardware and supporting anomaly injection for alert validation.
- **Dark-Mode Interface**: Designed for control room environments with a sleek, high-contrast UI using custom CSS.
- **Robust Data Handling**: Implements retry mechanisms and error logging for reliable IoT data fetching.

Technical Implementation

- **Frontend (Streamlit)**:
  - Built an interactive dashboard using Streamlit with a wide layout and auto-refresh (every 2 minutes) via `streamlit_autorefresh`.
  - Designed machine status cards with Plotly gauge charts for real-time vibration and temperature, using color-coded thresholds (green for normal, red for warning).
  - Created historical trend visualizations with Plotly line charts, showing vibration, temperature, and operating hours with threshold lines.
  - Applied custom CSS for a dark-mode interface, enhancing readability in low-light environments.
  - Managed chat history with `st.session_state` for a conversational AI experience.

- **IoT Integration**:
  - Used ThingSpeak's API to fetch real-time and historical sensor data (vibration, temperature, operating hours) for three machines (X, Y, Z).
  - Implemented `get_machine_data` and `get_historical_machine_data` with retry mechanisms (3 attempts, 5-second delays) to handle network issues.
  - Added a simulation mode (`simulate_sensor_data`) to generate synthetic data for testing, configurable via a sidebar button.

- **AI Engine**:
  - Integrated LlamaIndex with GroqLLM (Llama 3 70B) for natural language querying and maintenance recommendations.
  - Used `HuggingFaceEmbedding` (BAAI/bge-small-en-v1.5) to create embeddings for machine status documents, stored in a `VectorStoreIndex` for retrieval.
  - Built a `ReActAgent` with a custom system prompt defining operational thresholds (e.g., vibration <5 mm/s) and increased `max_iterations` to 10 for robust query handling.
  - Created dynamic `Document` objects from real-time sensor data, enabling the AI to answer queries like "Is Machine X due for maintenance?"

- **Core Technologies**:
    - **Streamlit**: Enabled rapid development of an interactive, real-time dashboard.
    - **LlamaIndex**: Simplified the creation of a queryable index for machine data, integrated with GroqLLM for AI responses.
    - **GroqLLM**: Powered natural language processing via the Groq API (Llama 3 70B model).
    - **Plotly**: Provided interactive, customizable visualizations for gauges and trend charts.
    - **ThingSpeak API**: Facilitated IoT data integration for real-time and historical insights.
    - **HuggingFace Embeddings**: Generated semantic embeddings for efficient data retrieval.

Challenges and Solutions

- **Challenge**: Learning LlamaIndex and integrating it with real-time IoT data as a self-taught developer.
    - **Solution**: Studied LlamaIndex documentation and experimented with `VectorStoreIndex` and `ReActAgent` to build a robust query engine.
- **Challenge**: Ensuring reliable data fetching from ThingSpeak's API under network instability.
    - **Solution**: Implemented a retry mechanism with exponential backoff and comprehensive error logging to handle failures gracefully.
- **Challenge**: Designing an intuitive dashboard for control room use.
    - **Solution**: Used custom CSS for a dark-mode UI and Plotly for interactive, color-coded visualizations that highlight critical thresholds.
- **Challenge**: Simulating realistic IoT data without hardware.
    - **Solution**: Developed a simulation mode that generates synthetic data with configurable parameters, allowing testing of alerts and AI responses.
- **Challenge**: Avoiding premature termination of AI query processing.
    - **Solution**: Increased the `max_iterations` parameter for the `ReActAgent` to ensure complex queries are fully processed.

Impact

The Manufacturing IoT Assistant empowers plant managers and engineers to monitor and maintain machinery efficiently, reducing downtime and improving operational decision-making. Key outcomes include:

- **Real-Time Insights**: The dashboard provides instant visibility into machine health, with alerts for high vibration, temperature, or maintenance needs.
- **Proactive Maintenance**: The AI assistant delivers actionable recommendations, such as "Lubricate Machine X due to 200+ operating hours," based on real-time data.
- **Accessibility**: Simulation mode enables testing and demonstration without IoT hardware, making the tool versatile for development and evaluation.
- **Skill Development**: Through self-learning, I mastered IoT integration, real-time dashboards, and AI-driven querying, enhancing my expertise in full-stack AI applications.
- **Portfolio Strength**: The project highlights my ability to build practical, industry-relevant tools, making it a compelling addition to my portfolio.

Lessons Learned

- **Real-Time Data Handling**: Implementing retries and error logging is critical for reliable IoT integrations.
- **User-Centric Design**: A dark-mode, high-contrast UI improves usability in operational environments like control rooms.
- **AI Optimization**: Tuning parameters like `max_iterations` and embedding models ensures robust AI performance.
- **Modular Architecture**: Separating data fetching, visualization, and AI querying into distinct functions improved code maintainability.
- **Self-Learning Resilience**: Overcoming the learning curve of LlamaIndex and IoT APIs reinforced the value of experimentation and persistence.

Future Enhancements

To productionize the application, I would:

- Integrate a database (e.g., PostgreSQL) to store historical sensor data and chat history, replacing in-memory `st.session_state`.
- Use a managed IoT platform (e.g., AWS IoT Core) for scalable, secure data ingestion from physical sensors.
- Implement real-time WebSocket updates for the dashboard instead of periodic refreshes to reduce latency.
- Add user authentication and role-based access to secure the application for multi-user environments.
- Optimize AI performance by exploring smaller embedding models or quantized LLMs for resource-constrained settings.

Screenshots:

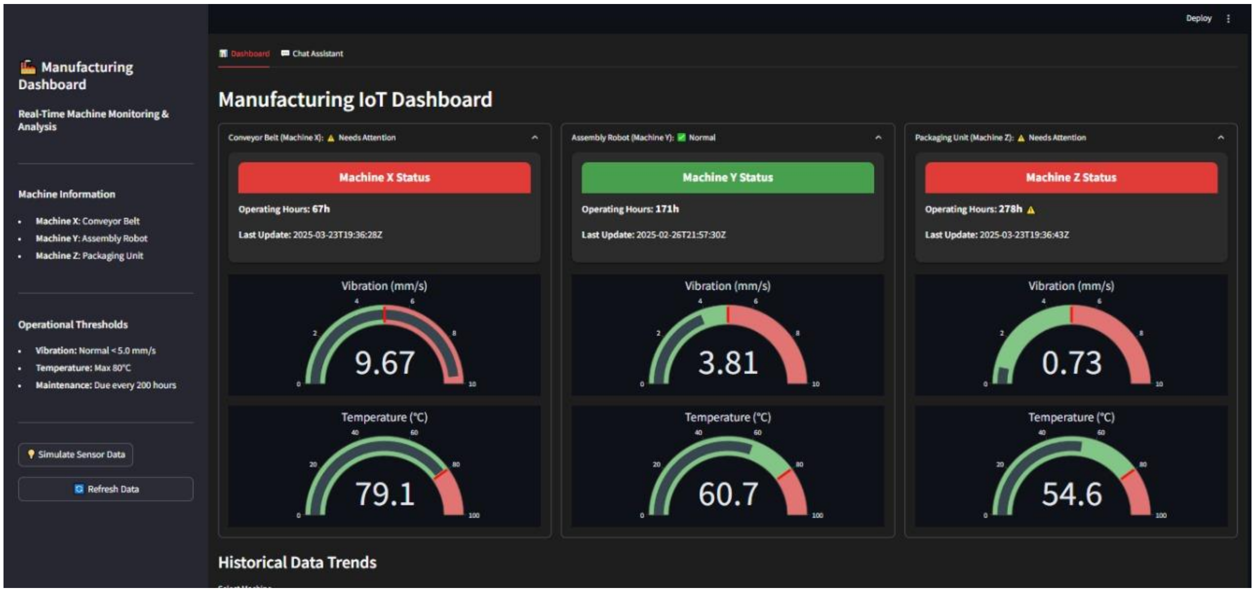

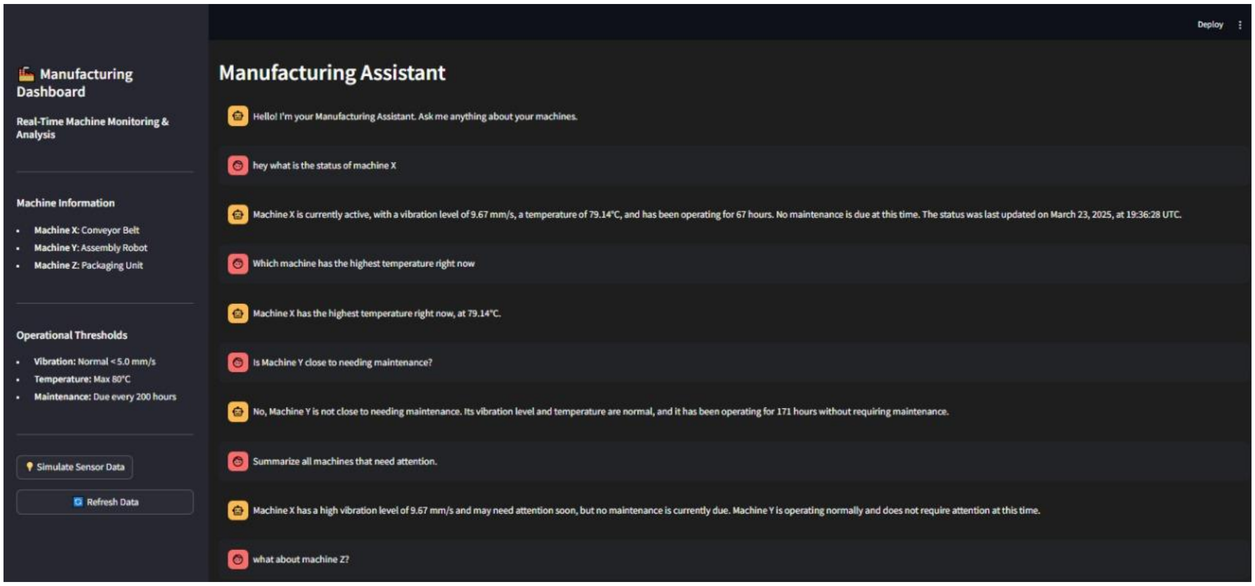
Deploy
Manufacturing Dashboard
Real-Time Machine Monitoring & Analysis
Machine Information
Machine X: Conveyor Belt
Machine Y: Assembly Robot
Machine Z: Packaging Unit
Operational Thresholds
Vibration: Normal < 5.0 mm/s
Temperature: Max 80°C
Maintenance: Due every 200 hours
Simulate Sensor Data
Refresh Data
Manufacturing Assistant
Hello! I'm your Manufacturing Assistant. Ask me anything about your machines.
hey what is the status of machine X
Machine X is currently active, with a vibration level of 9.67 mm/s, a temperature of 79.14°C, and has been operating for 67 hours. No maintenance is due at this time. The status was last updated on March 23, 2025, at 19:36:28 UTC.
Which machine has the highest temperature right now
Machine X has the highest temperature right now, at 79.14°C.
Is Machine Y close to needing maintenance?
No, Machine Y is not close to needing maintenance. Its vibration level and temperature are normal, and it has been operating for 171 hours without requiring maintenance.
Summarize all machines that need attention.
Machine X has a high vibration level of 9.67 mm/s and may need attention soon, but no maintenance is currently due. Machine Y is operating normally and does not require attention at this time.
what about machine Z?

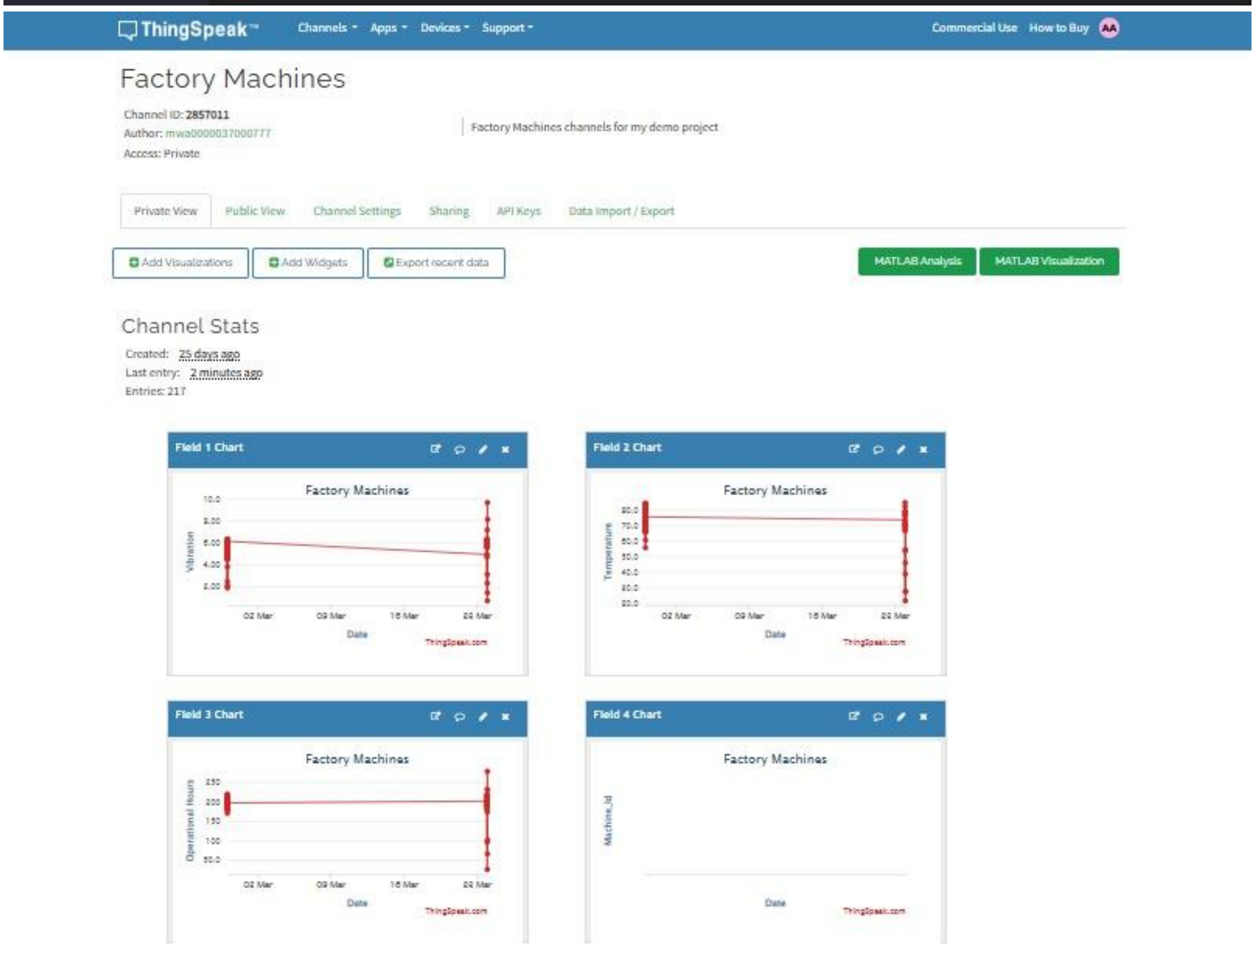
ThingSpeak™
Channels
Apps
Devices
Support
Commercial Use
How to Buy
Factory Machines
Channel ID: 2857011
Author: mwa0000037000777
Access: Private
Factory Machines channels for my demo project
Private View
Public View
Channel Settings
Sharing
API Keys
Data Import / Export
Add Visualizations
Add Widgets
Export recent data
MATLAB Analysis
MATLAB Visualization
Channel Stats
Created: 25 days ago
Last entry: 2 minutes ago
Entries: 217
Field 1 Chart
Factory Machines
Vibration
Date
ThingSpeak.com
Field 2 Chart
Factory Machines
Temperature
Date
ThingSpeak.com
Field 3 Chart
Factory Machines
Operational Hours
Date
ThingSpeak.com
Field 4 Chart
Factory Machines
Machine_Id
Date
ThingSpeak.com

```
Answer: Check the machine's balance, alignment, and perform a visual inspection of the par
ts to identify the cause of the high vibration. Additionally, review recent changes to the
 machine's operation and inspect the foundation to ensure it is secure.
INFO:root:Fetched data for Machine X: {'created_at': '2025-03-23T19:36:28Z', 'entry_id': 2
14, 'field1': '9.67', 'field2': '79.14', 'field3': '67', 'field4': 'X'}
INFO:root:Fetched data for Machine Y: {'created_at': '2025-02-26T21:57:30Z', 'entry_id': 1
85, 'field1': '3.81188195788643', 'field2': '60.72990164434485', 'field3': '171', 'field4'
: 'Y'}
INFO:root:Fetched data for Machine Z: {'created_at': '2025-03-23T19:36:43Z', 'entry_id': 2
15, 'field1': '0.73', 'field2': '54.58', 'field3': '278', 'field4': 'Z'}
INFO:root:Fetched data for Machine X: {'created_at': '2025-03-23T19:36:28Z', 'entry_id': 2
14, 'field1': '9.67', 'field2': '79.14', 'field3': '67', 'field4': 'X'}
INFO:root:Fetched data for Machine Y: {'created_at': '2025-02-26T21:57:30Z', 'entry_id': 1
85, 'field1': '3.81188195788643', 'field2': '60.72990164434485', 'field3': '171', 'field4'
: 'Y'}
INFO:root:Fetched data for Machine Z: {'created_at': '2025-03-23T19:36:43Z', 'entry_id': 2
15, 'field1': '0.73', 'field2': '54.58', 'field3': '278', 'field4': 'Z'}
INFO:root:Fetched data for Machine X: {'created_at': '2025-03-23T19:36:28Z', 'entry_id': 2
14, 'field1': '9.67', 'field2': '79.14', 'field3': '67', 'field4': 'X'}
INFO:root:Fetched data for Machine Y: {'created_at': '2025-02-26T21:57:30Z', 'entry_id': 1
85, 'field1': '3.81188195788643', 'field2': '60.72990164434485', 'field3': '171', 'field4'
: 'Y'}
INFO:root:Fetched data for Machine Z: {'created_at': '2025-03-23T19:36:43Z', 'entry_id': 2
15, 'field1': '0.73', 'field2': '54.58', 'field3': '278', 'field4': 'Z'}
Batches: 100%|██████████████████████████████| 1/1 [00:00<00:00, 11.07it/s]
> Running step 5c3b4744-2289-4d95-809e-b0ff5904d900. Step input: Do we need to replace any
 parts for Machine Z soon?
INFO:httpx:HTTP Request: POST https://api.groq.com/openai/v1/chat/completions "HTTP/1.1 20
```

Conclusion

The Manufacturing IoT Assistant is a testament to my self-learning journey and passion for building industry-relevant AI solutions. By integrating real-time IoT data with an AI-powered assistant, I created a tool that bridges the gap between raw sensor data and actionable insights. This project showcases my skills in IoT integration, real-time visualizations, and natural language processing, preparing me to tackle complex challenges in manufacturing and beyond.