

Case Study: Context-Aware Document Assistant

By Syed Armghan Ahmad – syedarmghanahmad.work@gmail.com

[Linkedin](#) - [GitHub](#) - [Repo for this project](#)

Overview

As a self-taught developer passionate about AI and natural language processing, I built the **Context-Aware Document Assistant**, a Retrieval-Augmented Generation (RAG) application that enables users to query large documents (PDFs or websites) using natural language. This project, developed through self-learning, integrates advanced AI tools like LangChain, HuggingFace, FAISS, and GroqLLM, with a Streamlit frontend for user interaction and a FastAPI backend for programmatic access. The application demonstrates my ability to design and implement a full-stack AI solution, addressing real-world challenges in document processing and information retrieval.

Problem Statement

In today's information-rich world, extracting specific insights from large documents, such as research papers, legal contracts, or websites, is time-consuming and error-prone. Traditional search methods often fail to provide context-aware answers, and users struggle to navigate multiple document types efficiently. My goal was to create a tool that:

- Allows users to ask questions in natural language and receive accurate, context-grounded responses.
- Supports both PDFs and website content with seamless context switching.
- Provides a user-friendly interface and a scalable API for diverse use cases.

Approach

To address this challenge, I designed a modular RAG-based system that combines document retrieval with AI-generated responses. The project consists of two main components: a Streamlit-based frontend for interactive use and a FastAPI-based backend for API access.

Key Features

- **Hybrid Document Contexts:** Users can upload PDFs or input website URLs, with the system maintaining separate vector stores for each context to ensure accurate retrieval.
- **Smart Document Processing:** Documents are split into configurable chunks (default: 1000 characters, 200 overlap), converted to embeddings using HuggingFace, and indexed in FAISS for efficient similarity search.
- **AI-Powered Responses:** Leverages GroqLLM models (e.g., Mixtral-8x7B) via LangChain to generate answers based on retrieved document chunks.
- **User-Friendly Interface:** Streamlit provides an intuitive UI with real-time chat, document upload, and settings for chunk size, overlap, and model selection.
- **Scalable API:** FastAPI endpoints enable programmatic document processing, querying, and chat history retrieval, suitable for integration into larger systems.
- **Error Handling and Feedback:** Robust error handling ensures graceful recovery from issues like invalid URLs or file uploads, with clear user feedback.

Technical Implementation

- **Frontend (Streamlit):**
 - Built an interactive UI using Streamlit, with a sidebar for configuring input type (PDFs or websites), chunking parameters, and model selection.
 - Used `session_state` to manage chat history, processed documents, and vector stores in-memory.
 - Implemented document processing with `PyPDFLoader` for PDFs and `WebBaseLoader` for websites, splitting text with `RecursiveCharacterTextSplitter`.
 - Integrated LangChain's RAG pipeline, combining FAISS-based retrieval with GroqLLM's generation capabilities.
- **Backend (FastAPI):**
 - Developed RESTful API endpoints (`/process-documents`, `/query`, `/chat-history`, `/health`) using FastAPI.
 - Used Pydantic models for request validation and a global state dictionary for simplicity (noted for replacement with a database in production).
 - Mirrored the frontend's RAG pipeline, ensuring consistency in document processing and query handling.
- **Core Technologies:**
 - **LangChain:** Simplified the RAG pipeline with modular chains (`create_stuff_documents_chain`, `create_retrieval_chain`).
 - **HuggingFace Embeddings:** Generated semantic embeddings for document chunks, enabling accurate retrieval.
 - **FAISS:** Provided efficient vector storage and similarity search for large document sets.
 - **GroqLLM:** Powered context-aware responses via the Groq API.
 - **Streamlit/FastAPI:** Enabled rapid prototyping (Streamlit) and scalable API development (FastAPI).

Challenges and Solutions

- **Challenge:** Learning and integrating complex AI tools like LangChain and FAISS as a self-taught developer.
 - **Solution:** Leveraged official documentation, tutorials, and experimentation to understand RAG pipelines, embeddings, and vector stores.
- **Challenge:** Managing multiple document types without mixing contexts.
 - **Solution:** Designed separate FAISS indices (`pdf_vector`, `web_vector`) and a `current_context` tracker to ensure context-specific retrieval.
- **Challenge:** Optimizing document chunking for retrieval accuracy.
 - **Solution:** Experimented with chunk sizes and overlaps, settling on defaults that balanced context retention and performance.
- **Challenge:** Ensuring robust error handling for user inputs.
 - **Solution:** Implemented try-except blocks and user feedback (e.g., `st.error`, `HTTPException`) to handle invalid URLs, file uploads, or API issues.

Impact

The Context-Aware Document Assistant successfully enables users to query large documents efficiently, delivering accurate answers with source attribution (e.g., PDF page numbers or website URLs). Key outcomes include:

- **Enhanced Productivity:** Users can extract insights from documents without manual searching, saving time in scenarios like research or legal analysis.
- **Flexibility:** Supports diverse document types and customizable settings, making it adaptable to various use cases.
- **Skill Development:** Through self-learning, I gained expertise in RAG, vector embeddings, full-stack development, and API design, preparing me for real-world AI projects.
- **Portfolio Strength:** The project showcases my ability to build end-to-end AI applications, from front-end design to back-end scalability, making it a standout piece in my portfolio.

Lessons Learned

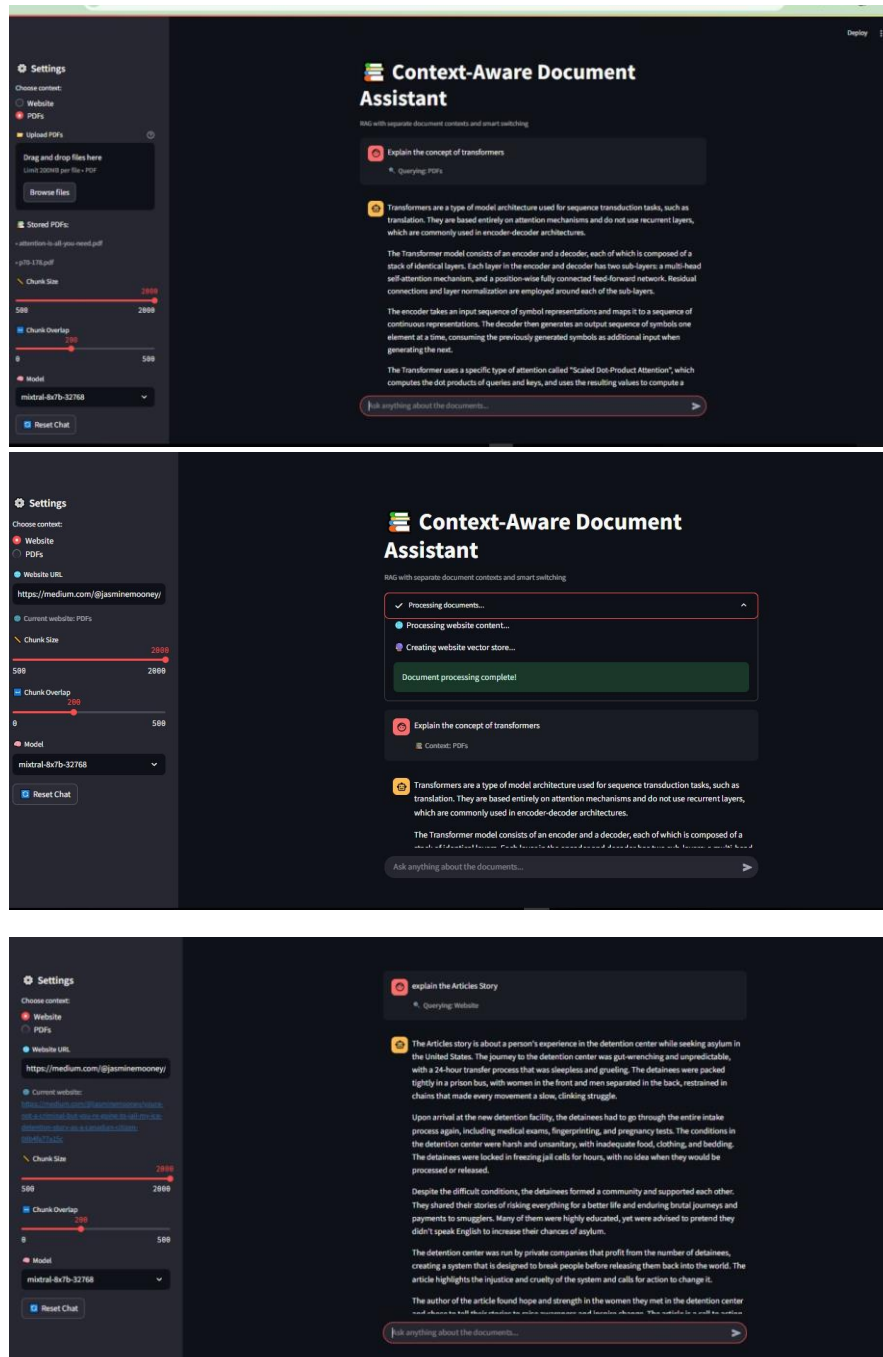
- **Modular Design:** Structuring the code into reusable functions (e.g., `process_documents`, `create_chain`) improved maintainability and extensibility.
- **Iterative Learning:** Experimentation with chunking parameters and model selection taught me the importance of tuning AI systems for specific tasks.
- **User-Centric Approach:** Prioritizing clear feedback and intuitive UI design enhanced the user experience.
- **Production Readiness:** In-memory state management works for prototyping but requires a database for scalability, informing future projects.

Future Enhancements

To productionize the application, I would:

- Replace in-memory state with a database (e.g., PostgreSQL) for persistent storage of chat history and document metadata.
- Use a managed vector database (e.g., Pinecone) for scalable, distributed vector storage.
- Implement async document processing with Celery and Redis to handle large document sets.
- Add user authentication and rate limiting to secure the FastAPI backend.
- Explore smaller models (e.g., DistilBERT) for resource-constrained environments.

Screenshots:



Conclusion

The Context-Aware Document Assistant is a testament to my self-learning journey and passion for AI-driven solutions. By tackling a real-world problem with a sophisticated tech stack, I demonstrated my ability to design, implement, and iterate on a full-stack AI application. This project not only solved a practical challenge but also equipped me with skills in RAG, vector search, and web development, preparing me for impactful contributions in professional settings.